

Short answer: don't use a general text embedding model on raw FEN. You'll get token-level artifacts that don't reflect chess similarity. Use a purpose-built board encoder and (optionally) learn a small embedding head.

Here's a practical, RAG-friendly stack that works well for chess positions:

1) Canonicalization & dedup (pre-embedding)

- **Normalize FEN:** drop halfmove/fullmove counters; keep side-to-move, castling, en-passant.
- **Symmetry canonicalization:** pick a canonical transform over {identity, flip ranks/files, rotate 180°} that makes equivalent patterns map to the same representative (e.g., always put the side-to-move as "White", flip board if needed).
- **Hash:** store a Zobrist (or any 64-bit) hash of the canonical FEN for exact dedup and a unique key.

2) Baseline embedding (deterministic, zero-training)

Bitboard one-hot → float vector

- 12 piece planes × 64 squares = **768 dims** (binary 0/1).
- ◦ side-to-move (1), castling rights (4 one-hots), en-passant file (9 one-hots incl. "none") → **~782 dims total**.
- Store as float32 and use **cosine** distance.
This already gives sensible "same/similar piece layouts ≈ close" behavior.

3) Motif-aware head (small, optionally trained)

If you care about *strategic* motifs (e.g., queenside majority), add a compact feature block and/or learn a projection:

- **Pawn-structure features (~40–60 dims)**
 - File counts: white pawns fA...fH (8), black pawns fA...fH (8).

- Binary flags: isolated, doubled, backward, passed (per file, compressed counts).
- Pawn-island count (white, black), open/semi-open files per side, majority indicators (queenside/kingside).
- **Material & king-safety lite (~10–20 dims)**
 - Material vector (piece counts per type), king files/ranks bucketed, castling availability.
- **Concatenate** with the 782-d baseline and (optionally) learn a **64–128-d projection** with a tiny MLP trained contrastively:
 - **Positives:** positions equal under symmetries; same pawn structure; small Stockfish eval delta; same ECO bucket in early middlegame.
 - **Negatives:** random positions; different pawn structures.

(Unverified) In practice, this small learned head often outperforms raw one-hot for motif queries while staying cheap to compute.

4) Engine-informed enrichment (optional)

If you can batch-evaluate positions:

- Store **eval score**, **phase** (opening/middlegame/endgame), and a few **engine features** (e.g., mobility, passed pawns).
- Use them as **filters** or as extra dimensions (scaled); don't force them into the core embedding if you worry about engine bias.

5) Vector-DB setup

- Use **cosine** or **dot-product** for dense floats.
- Keep two indexes:
 1. **layout index** = the 782-d (or projected 64–128-d) vector.
 2. **pawn-structure index** = a 32–64-d pawn-only vector (fast shortlist for structure-first searches).

- Store the **canonical FEN** + metadata (game id, move number, ECO, eval) alongside vectors.

6) Query patterns that work

- “**Queenside majority attack**”: search the **pawn-structure index** first (queenside majority flag + file counts), then re-rank with the layout index or engine eval proximity.
- “**Same piece placement / tactical pattern**”: hit the **layout index** directly.
- **Filters**: side-to-move, castling rights, eval window, phase, ECO.

Why not a text embedding (OpenAI, bge, etc.) on FEN?

- FEN is a compact code, not natural language; tokenizers will slice it arbitrarily.
- Distance in token space doesn't track **board** similarity.

(Unverified) You may see occasional wins, but it's brittle and hard to reason about.

Minimal implementation sketch (what to build first)

1. **Canonicalizer** (FEN ↔ board; symmetry normalize; strip clocks).
2. **Encoder**:
 - `float[782]` baseline vector (12×64 pieces + STM + castling + EP).
 - `float[?]` pawn-features vector (counts/flags).
3. **(Optional)** Tiny MLP to 128-d with contrastive training (you can start untrained and add this later).
4. **Two indexes** in your vector DB, plus scalar metadata columns for fast filtering.

If you want, I can generate a reference encoder (Python or OCaml) that outputs both the 782-d baseline and a 48-d pawn-structure vector, plus a Qdrant schema to store them.

Below is a self-contained OCaml reference encoder for chess positions in FEN.

It produces:

- a **782-d baseline vector**: 12×64 one-hot piece planes (P,N,B,R,Q,K × white/black = 768) + side-to-move (1) + castling rights K,Q,k,q (4) + en-passant file one-hot (a–h, none = 9)
- a **48-d pawn/metadata vector** with file counts, islands, open/semi-open files, isolated/doubled/passed counts, side majorities, material counts, castling/STM/EP flags, and a crude “phase” scalar

No external deps. Keep or adapt the feature ordering as you need.

```
(* File: chess_fen_encoder.ml *)

module ChessFenEncoder = struct
  (*
    --- Public API -----
    -----

    val encode_layout    : string -> float array    (* 782-d
baseline *)
    val encode_features  : string -> float array    (* 48-d
pawn+meta *)
    val encode_all       : string -> float array * float array

    -----
    -----

    Baseline (782 dims), in order:
      [0..767]    12 piece planes × 64 squares (a1..h1, a2..h2,
... a8..h8)
                                Plane order: WP, WN, WB, WR, WQ, WK, BP, BN,
BB, BR, BQ, BK
```

[768] side-to-move: 1.0 = White, 0.0 = Black
[769..772] castling rights one-hot flags: K, Q, k, q (1.0 or 0.0)
[773..781] en-passant file one-hot: a..h plus index 8 = "none"

Features (48 dims), in order:

[0.. 7] white pawn file counts a..h
[8..15] black pawn file counts a..h
[16] white pawn islands
[17] black pawn islands
[18] open files (no pawns either side)
[19] semi-open for White (no white pawns, ≥ 1 black pawn)
[20] semi-open for Black (no black pawns, ≥ 1 white pawn)
[21] isolated white pawns
[22] isolated black pawns
[23] doubled white pawns (sum over files $\max(0, \text{count}-1)$)
[24] doubled black pawns
[25] passed white pawns (simplified, see code)
[26] passed black pawns (simplified)
[27] queenside majority White (a-d > e-h) {0.0,1.0}
[28] queenside majority Black (a-d > e-h) {0.0,1.0}
[29] kingside majority White (e-h > a-d) {0.0,1.0}
[30] kingside majority Black (e-h > a-d) {0.0,1.0}
[31..36] white material counts: P,N,B,R,Q,K
[37..42] black material counts: P,N,B,R,Q,K
[43] castling right exists (White): {0.0,1.0}
[44] castling right exists (Black): {0.0,1.0}
[45] side-to-move (1.0 white, 0.0 black)
[46] en-passant available (not "-"): {0.0,1.0}
[47] phase (crude 0.0 opening .. 1.0 endgame; see code)

Notes:

- Coordinates: squares are indexed a1..h1, a2..h2, ... , a8..h8 (file-major).

- FEN parsing ignores halfmove/fullmove counters in the baseline planes.

*)

(* --- Types & basics -----
----- *)

type color = White | Black

type piece = Pawn | Knight | Bishop | Rook | Queen | King

let is_between x a b = (x >= a) && (x <= b)

let piece_of_char = function

| 'P' -> Some (White, Pawn)
| 'N' -> Some (White, Knight)
| 'B' -> Some (White, Bishop)
| 'R' -> Some (White, Rook)
| 'Q' -> Some (White, Queen)
| 'K' -> Some (White, King)
| 'p' -> Some (Black, Pawn)
| 'n' -> Some (Black, Knight)
| 'b' -> Some (Black, Bishop)
| 'r' -> Some (Black, Rook)
| 'q' -> Some (Black, Queen)
| 'k' -> Some (Black, King)
| _ -> None

let plane_index (c, p) =

match c, p with

| White, Pawn -> 0
| White, Knight -> 1
| White, Bishop -> 2
| White, Rook -> 3
| White, Queen -> 4
| White, King -> 5
| Black, Pawn -> 6
| Black, Knight -> 7
| Black, Bishop -> 8
| Black, Rook -> 9
| Black, Queen -> 10

```
| Black, King    -> 11
```

```
(* board: 64 squares, each is (color * piece) option
   index: file 0..7 (a..h), rank 0..7 (1..8); idx = file +
8*rank *)
type board = (color * piece) option array

type castling = {
  wk : bool;   (* K *)
  wq : bool;   (* Q *)
  bk : bool;   (* k *)
  bq : bool;   (* q *)
}

type fen = {
  board : board;
  stm    : color;           (* side to move *)
  cast   : castling;
  ep     : int option;      (* en-passant file 0..7, or None
*)
  (* clocks ignored for encoding purposes *)
}

(* --- FEN parsing -----
----- *)

let empty_board () : board = Array.make 64 None

let idx_of_file_rank file rank = (file land 7) + (rank land
7) * 8

let parse_ep (s : string) : int option =
  (* s is like "e3" or "-" *)
  if s = "-" then None
  else
    let file_ch = s.[0] in
    if is_between (Char.code file_ch) (Char.code 'a')
(Char.code 'h')
    then Some (Char.code file_ch - Char.code 'a')
    else None

let parse_castling (s : string) : castling =
```

```

if s = "-" then {wk=false; wq=false; bk=false; bq=false}
else
  let has c = String.exists (fun x -> x = c) s in
  {
    wk = has 'K';
    wq = has 'Q';
    bk = has 'k';
    bq = has 'q';
  }

let parse_fen (fen_str : string) : fen =
  (* FEN: pieces / stm / castling / ep / halfmove / fullmove
  *)
  let parts = String.split_on_char ' ' fen_str in
  match parts with
  | pieces :: stm :: cast :: ep :: _half :: _full :: _ ->
    let b = empty_board () in
    (* pieces rows: rank 8 to 1, slash-separated *)
    let ranks = String.split_on_char '/' pieces in
    if List.length ranks <> 8 then failwith "Invalid FEN:
ranks";
    List.iteri
      (fun i rstr ->
        (* i=0 -> rank 8; want to place on rank=7 down to
0 *)
        let rank = 7 - i in
        let file = ref 0 in
        String.iter (fun ch ->
          if is_between (Char.code ch) (Char.code '1')
(Char.code '8')
          then file := !file + (Char.code ch - Char.code
'0')
          else match piece_of_char ch with
            | Some cp ->
              if !file >= 8 then failwith "Invalid FEN:
file overflow";
              let idx = idx_of_file_rank !file rank in
              b.(idx) <- Some cp;
              incr file
            | None -> failwith "Invalid FEN: piece char"
          ) rstr;

```



```

        if !file <> 8 then failwith "Invalid FEN: file
count"

    ) ranks;
    let side =
        match stm with
        | "w" -> White
        | "b" -> Black
        | _ -> failwith "Invalid FEN: stm"
    in
    let cast = parse_castling cast in
    let epf = parse_ep ep in
    { board = b; stm = side; cast; ep = epf }
    | _ -> failwith "Invalid FEN: fields"

(* --- Helpers for features -----
----- *)

let is_pawn = function (_, Pawn) -> true | _ -> false
let is_white = function White, _ -> true | _ -> false
let is_black = function Black, _ -> true | _ -> false

let piece_value = function
    | Pawn -> 1 | Knight -> 3 | Bishop -> 3 | Rook -> 5 | Queen
-> 9 | King -> 0

let material_counts (b:board) =
    (* counts: White P,N,B,R,Q,K ; Black P,N,B,R,Q,K *)
    let w = Array.make 6 0 and bl = Array.make 6 0 in
    Array.iter (function
        | None -> ()
        | Some (c,p) ->
            let i = match p with
                | Pawn -> 0 | Knight -> 1 | Bishop -> 2
                | Rook -> 3 | Queen -> 4 | King -> 5
            in
            (match c with
                | White -> w.(i) <- w.(i) + 1
                | Black -> bl.(i) <- bl.(i) + 1)
    ) b;
    (w, bl)

let file_of_idx idx = idx mod 8

```

```

let rank_of_idx idx = idx / 8

let pawn_file_counts (b:board) =
  let wf = Array.make 8 0 and bf = Array.make 8 0 in
  Array.iteri (fun idx cell ->
    match cell with
    | Some (White, Pawn) -> wf.(file_of_idx idx) <- wf.
(file_of_idx idx) + 1
    | Some (Black, Pawn) -> bf.(file_of_idx idx) <- bf.
(file_of_idx idx) + 1
    | _ -> ()
  ) b;
  (wf, bf)

let islands_of_files (fcounts:int array) =
  (* number of contiguous groups of files with count>0 *)
  let islands = ref 0 in
  let in_group = ref false in
  for f=0 to 7 do
    if fcounts.(f) > 0 then (
      if not !in_group then (in_group := true; incr islands)
    ) else in_group := false
  done;
  !islands

let open_and_semiopen (wf:int array) (bf:int array) =
  let open_files = ref 0
  and semi_w = ref 0
  and semi_b = ref 0 in
  for f=0 to 7 do
    let w = wf.(f) and b = bf.(f) in
    if w = 0 && b = 0 then incr open_files;
    if w = 0 && b > 0 then incr semi_w;
    if b = 0 && w > 0 then incr semi_b;
  done;
  (!open_files, !semi_w, !semi_b)

let isolated_pawns (fcounts:int array) =
  let iso = ref 0 in
  for f=0 to 7 do
    if fcounts.(f) > 0 then
      let left = if f>0 then fcounts.(f-1) else 0 in

```

```

        let right = if f<7 then fcounts.(f+1) else 0 in
        if left = 0 && right = 0 then iso := !iso + fcounts.(f)
done;
!iso

let doubled_pawns (fcounts:int array) =
  let d = ref 0 in
  for f=0 to 7 do
    if fcounts.(f) > 1 then d := !d + (fcounts.(f) - 1)
  done;
  !d

(* Passed pawns (simplified):
   - White pawn on (f,r) is passed if there is NO black pawn
on files f-1,f,f+1
   on ranks strictly greater than r (ahead of it).
   - Symmetric for Black (look for white pawns behind).
*)
let passed_pawns (b:board) =
  let w_passed = ref 0 and b_passed = ref 0 in
  let black_pawns_on f r =
    (* is there a black pawn on files f-1..f+1 with rank>r ?
    *)
    let found = ref false in
    for ff = max 0 (f-1) to min 7 (f+1) do
      for rr = r+1 to 7 do
        let idx = idx_of_file_rank ff rr in
        match b.(idx) with
        | Some (Black, Pawn) -> found := true
        | _ -> ()
      done
    done;
    !found
  in
  let white_pawns_on f r =
    (* is there a white pawn on files f-1..f+1 with rank<r ?
    *)
    let found = ref false in
    for ff = max 0 (f-1) to min 7 (f+1) do
      for rr = r-1 downto 0 do
        let idx = idx_of_file_rank ff rr in
        match b.(idx) with

```

```

        | Some (White, Pawn) -> found := true
        | _ -> ()
    done
done;
!found
in
Array.iteri (fun idx cell ->
    match cell with
    | Some (White, Pawn) ->
        let f = file_of_idx idx and r = rank_of_idx idx in
        if not (black_pawns_on f r) then incr w_passed
    | Some (Black, Pawn) ->
        let f = file_of_idx idx and r = rank_of_idx idx in
        if not (white_pawns_on f r) then incr b_passed
    | _ -> ()
) b;
(!w_passed, !b_passed)

let majority_flags (wf:int array) (bf:int array) =
    let sum range arr =
        let s = ref 0 in
        List.iter (fun f -> s := !s + arr.(f)) range; !s
    in
    let w_q = sum [0;1;2;3] wf
    and w_k = sum [4;5;6;7] wf
    and b_q = sum [0;1;2;3] bf
    and b_k = sum [4;5;6;7] bf in
    let w_q_maj = if w_q > w_k then 1.0 else 0.0
    and b_q_maj = if b_q > b_k then 1.0 else 0.0
    and w_k_maj = if w_k > w_q then 1.0 else 0.0
    and b_k_maj = if b_k > b_q then 1.0 else 0.0 in
    (w_q_maj, b_q_maj, w_k_maj, b_k_maj)

let phase_scalar (w:int array) (bl:int array) =
    (* crude phase: normalize non-pawn material to [0..1]
endgame-ish *)
    let nonpawn_sum counts =
        let n = counts.(1) and b = counts.(2) and r = counts.(3)
    and q = counts.(4) in
        (n * 3) + (b * 3) + (r * 5) + (q * 9)
    in
    let tot = nonpawn_sum w + nonpawn_sum bl in

```

```

    let max_tot = 62. (* start of game: N4,B4,R4,Q2 ->
12+12+20+18 = 62 *)
    let x = float_of_int tot /. max_tot in
    let p = 1.0 -. x in
    if p < 0.0 then 0.0 else if p > 1.0 then 1.0 else p

(* --- Baseline encoder (782-d) -----
----- *)

let encode_layout (fen_str : string) : float array =
  let f = parse_fen fen_str in
  let v = Array.make 782 0.0 in
  (* 12 planes × 64 squares *)
  Array.iteri (fun idx cell ->
    match cell with
    | None -> ()
    | Some cp ->
      let pidx = plane_index cp in
      let base = pidx * 64 in
      v.(base + idx) <- 1.0
  ) f.board;
  (* side to move *)
  v.(768) <- (match f.stm with White -> 1.0 | Black -> 0.0);
  (* castling K,Q,k,q *)
  v.(769) <- if f.cast.wk then 1.0 else 0.0;
  v.(770) <- if f.cast.wq then 1.0 else 0.0;
  v.(771) <- if f.cast.bk then 1.0 else 0.0;
  v.(772) <- if f.cast.bq then 1.0 else 0.0;
  (* en-passant file one-hot a..h + none *)
  let ep_idx = match f.ep with Some file -> file | None -> 8
in
  for i=0 to 8 do v.(773 + i) <- if i = ep_idx then 1.0 else
0.0 done;
  v

(* --- Feature encoder (48-d) -----
----- *)

let encode_features (fen_str : string) : float array =
  let f = parse_fen fen_str in
  let feats = Array.make 48 0.0 in
  let wf, bf = pawn_file_counts f.board in

```

```

(* file counts *)
for i=0 to 7 do
  feats.(i)      <- float_of_int wf.(i);
  feats.(8 + i)  <- float_of_int bf.(i);
done;
(* islands *)
feats.(16) <- float_of_int (islands_of_files wf);
feats.(17) <- float_of_int (islands_of_files bf);
(* open/semi-open *)
let open_files, semi_w, semi_b = open_and_semiopen wf bf in
feats.(18) <- float_of_int open_files;
feats.(19) <- float_of_int semi_w;
feats.(20) <- float_of_int semi_b;
(* isolated/doubled *)
feats.(21) <- float_of_int (isolated_pawns wf);
feats.(22) <- float_of_int (isolated_pawns bf);
feats.(23) <- float_of_int (doubled_pawns wf);
feats.(24) <- float_of_int (doubled_pawns bf);
(* passed (simplified) *)
let w_passed, b_passed = passed_pawns f.board in
feats.(25) <- float_of_int w_passed;
feats.(26) <- float_of_int b_passed;
(* majorities *)
let wq, bq, wk, bk = majority_flags wf bf in
feats.(27) <- wq; feats.(28) <- bq; feats.(29) <- wk;
feats.(30) <- bk;
(* material counts *)
let wmat, bmat = material_counts f.board in
for i=0 to 5 do feats.(31 + i) <- float_of_int wmat.(i)
done;
for i=0 to 5 do feats.(37 + i) <- float_of_int bmat.(i)
done;
(* castling availability (any side) *)
feats.(43) <- if f.cast.wk || f.cast.wq then 1.0 else 0.0;
feats.(44) <- if f.cast.bk || f.cast.bq then 1.0 else 0.0;
(* stm / ep flag *)
feats.(45) <- (match f.stm with White -> 1.0 | Black ->
0.0);
feats.(46) <- (match f.ep with Some _ -> 1.0 | None ->
0.0);
(* phase *)
feats.(47) <- phase_scalar wmat bmat;

```

feats

```
let encode_all fen_str = (encode_layout fen_str,  
encode_features fen_str)
```

```
(* --- Pretty helpers (optional) -----  
----- *)
```

```
let pp_vec_short name (v:float array) n =  
  Printf.printf "%s (%d): [" name (Array.length v);  
  for i=0 to min (n-1) (Array.length v - 1) do  
    if i>0 then print_string ", ";  
    Printf.printf "%.0f" v.(i)  
  done;  
  if Array.length v > n then print_string ", ...";  
  print_endline "]"
```

```
end
```

```
(* ----- Quick usage example -----  
-----  
  let fen = "r4rk1/ppp1bppp/2n2B2/8/8/2P2P2/PPP1bP1P/1K1R2R1 w  
- - 0 17" in  
  let layout = ChessFenEncoder.encode_layout fen in  
  let feats = ChessFenEncoder.encode_features fen in  
  ChessFenEncoder.pp_vec_short "layout" layout 20;  
  ChessFenEncoder.pp_vec_short "feats" feats 48;  
-----  
----- *)
```